



Module_3_Project_2

Overview



LearningSteps Evolution — turning a manually-built Azure 2-tier app (Python FastAPI + PostgreSQL) into a fully automated DevSecOps pipeline: Docker → Terraform → AKS → GitHub Actions (with security gates) → optional Prometheus/Grafana monitoring.


Timeline: 4 weeks. Working environment: Kali Linux, project root `~/projects/Module_3_Project_2/learningsteps-project`.

This page is a living tracker — update the **Progress Tracker** section and the phase checkboxes as work advances.



Target Architecture

User traffic flows through a public load balancer into AKS, which reads its database secret from Key Vault and talks to PostgreSQL. Terraform owns every piece of infrastructure below.



Component	Purpose
VNet + subnets	Separate address space for AKS nodes and the database
AKS cluster	Runs the LearningSteps pods, self-heals and scales
PostgreSQL Flexible Server	Private-network database, stores journal records
Azure Container Registry	Stores built and scanned Docker images
Key Vault	Stores the DB connection string, delivered to pods via CSI driver
Public Load Balancer	Exposes the API to the internet via a Kubernetes Service



Tools & Installation (Kali Linux)

Install order matters a little (Azure CLI and Terraform each need a workaround on Kali — see Troubleshooting section).



Tool	Purpose	Install command
Docker	Build/run/test the container image	<code>sudo apt install docker.io -y</code>
Git	Version control, triggers the pipeline	<code>sudo apt install git -y</code>

Tool	Purpose	Install command
Python 3	Run the FastAPI app locally	<code>sudo apt install python3 python3-venv python3-pip -y</code>
Azure CLI	Create/manage Azure resources	<code>pipx install azure-cli</code>
Terraform	Provision infrastructure as code	<code>sudo apt install terraform -y</code> (after adding HashiCorp repo)
kubectl	Talk to the AKS cluster	<code>sudo az aks install-cli</code>
Helm	Deploy the app via the chart	<code>curl .../get-helm-3 bash</code>
Trivy	Image + Terraform vulnerability scanning	<code>curl .../install.sh sudo sh -s -- -b /usr/local/bin</code>
trufflehog	Secret scanning	<code>curl .../install.sh sh -s -- -b /usr/local/bin</code>

Full idempotent installer: attached at the bottom of this page (`install-tools-kali`).

Verify everything at once:

```
docker --version && git --version && python3 --version && az --version | head -1 && terraform --version | head -1 && kubectl version --client && helm version --short && trivy --version | head -1 && trufflehog --version
```



Step-by-Step Plan



Week	Focus	Tasks	Done when
1	Working app, then containerize	Fix CRUD endpoints locally, write Dockerfile, test with docker compose, push to ACR	App works, image runs from ACR
2	Infrastructure as code	Terraform for VNet, AKS, PostgreSQL, ACR, Key Vault	<code>terraform apply</code> builds everything from scratch
3	Connect + orchestrate + automate	Key Vault secret into app, K8s manifests/Helm, CI/CD with OIDC + security gates	Push to main deploys automatically, no secrets in repo
4	Prove it works	Autoscaling test, break the pipeline on purpose, optional monitoring, README	All four handover checks pass



Progress Tracker



Update this table as each phase completes. Last updated: current session.

Step	Status	Notes
0 — Plan & architecture	✅ Done	Reviewed project docs, mapped flow, built build guide
1 — Install tools on Kali	✅ Done	Re-verified after restart: all 9 tools, Azure login active, Docker works without sudo (hello-world OK)
2 — Get the app working locally	✅ Done	All 6 endpoints verified working (200/200/404 GET → DELETE → GET sequence confirmed after re-patching the router)
3 — Dockerize + push to ACR	✅ Done	Image built, scanned (44 HIGH / 0 CRITICAL), pushed to ACR (RG module3-project2, ACR m3p2acr, tag v1)
4 — Terraform infra	✅ Done	All 7 resources live: RG, VNet, 2 subnets, AKS (2 nodes, Standard_D2s_v6, OIDC issuer enabled), Postgres Flexible Server (zone 1), Private DNS zone+link, ACR (imported), Key Vault. Trivy config: 6 → 3 findings fixed (RBAC, network policy, KV ACLs/purge protection)

Step	Status	Notes
5 — Connect Key Vault secrets	✅ Done	database-url secret delivered via CSI + SecretProviderClass, verified end-to-end: pod reads DATABASE_URL, POST /entries succeeds against Azure Postgres
6 — Deploy with Helm chart	✅ Done	Chart created from scratch on Kali (was never unzipped), raw-YAML from Phase 5 replaced cleanly via helm upgrade --install. Self-healing, HPA, and public endpoint all verified; POST /entries returns 200 live
7 — CI/CD pipeline (OIDC + gates)	✅ Done	Full pipeline green: secrets scan → lint/test → build → Trivy scan → OIDC login → ACR push → Helm deploy. New pod live, public endpoint confirmed 200. Still open: prove the gate blocks a bad push
8 — Optional monitoring	⏸ Paused	Infra fully working (app instrumented, Prometheus scraping confirmed up, Grafana login OK); dashboard panels not finished — paused by decision since this step is optional, moving to Step 9
9 — Final demo	🔄 Next	

Phase 2 — Get the App Working

Goal: correct, fully working LearningSteps API running locally, before Docker touches it.



Corrected after inspecting the real repo: `requirements.txt` lives at `api/requirements.txt` (not root), the database name is `learning_journal` (not `learningsteps`), and the app must run from inside `api/` since `main.py` does relative imports like `from routers.journal_router import ...`.



Command	What it does
<code>git clone https://github.com/CyberstepsDE/learningsteps.git</code>	Downloads the app source
<code>python3 -m venv venv && source venv/bin/activate</code>	Creates and activates an isolated Python environment
<code>pip install -r api/requirements.txt</code>	Installs app dependencies (path is under <code>api/</code> , not root)
<code>cp .env-sample .env</code>	Creates the local env file (git-ignored); DATABASE_URL points at the <code>postgres</code> service/host
<code>docker run --name learningsteps-db -e POSTGRES_USER=postgres -e POSTGRES_PASSWORD=postgres -e POSTGRES_DB=learning_journal -p 5433:5432 -v "\$PWD/database_setup.sql":/docker-entrypoint-initdb.d/database_setup.sql:ro -d postgres:16</code>	Starts local Postgres with the real DB name and schema auto-loaded
<code>cd api && uvicorn main:app --reload --host 0.0.0.0 --port 8000</code>	Runs the API locally — must run from inside <code>api/</code>
<code>curl -X POST http://localhost:8000/entries -d '{"work":"...","struggle":"...","intention":"..."}'</code>	Tests creating a journal entry (real fields: work / struggle / intention)



Checkpoint result: 4 of 6 endpoints work correctly (POST /entries, GET /entries, DELETE /entries, PATCH /entries/{id}). Two are intentional `501 Not Implemented` stubs in the exercise code — `GET /entries/{id}` and `DELETE /entries/{id}`. Left unmodified for now by decision — not fixed yet. Phase closed as-is to move forward.



Repair update — stub endpoints fixed

`GET /entries/{entry_id}` and `DELETE /entries/{entry_id}` were left as `501 Not Implemented` stubs when Phase 2 first closed. Fixed before moving to Phase 3 so the container build starts from a fully working app.

GET /entries/{entry_id} — now calls `entry_service.get_entry()` (already implemented in the service layer, just never wired to the router) and returns 404 if not found, 200 with the entry otherwise.

DELETE /entries/{entry_id} — now checks existence via `entry_service.get_entry()`, returns 404 if missing, otherwise calls `entry_service.delete_entry()` and confirms deletion.
Verified: create → GET by id (200) → DELETE by id (200) → GET by id again (404).
All 6 endpoints now functional. Status:  Resolved.

Diagnosis — locating the stub code

Before patching, the file's actual state was confirmed (the first repair attempt silently failed to save — caught only because the verification step was re-run rather than trusted):

```
cd ~/projects/Module_3_Project_2/learningsteps-project/learningsteps
grep -n "501\|def get_entry\|def delete_entry" api/routers/journal_router.py
```

Greps the router file for the 501 status code and both function definitions, to confirm the exact current state before editing.

Verification — before the fix

```
BASE=http://localhost:8000

curl -s -X POST $BASE/entries -H "Content-Type: application/json" \
  -d '{"work": "Fix stub endpoints", "struggle": "None", "intention": "Move to Docker"}'
```

Creates a test entry and captures its `id` from the response, used as the target for the checks below.

```
ID=<paste-id-here>

echo "GET by id:"
curl -s -o /dev/null -w "%{http_code}\n" $BASE/entries/$ID

echo "DELETE by id:"
curl -s -o /dev/null -w "%{http_code}\n" -X DELETE $BASE/entries/$ID

echo "GET by id again (should be gone now):"
curl -s -o /dev/null -w "%{http_code}\n" $BASE/entries/$ID
```

Checks all three status codes in sequence. First run returned `501 / 501 / 501` — confirming the stubs were still live despite the edit appearing complete.

Repair — patch script used

Applied with an exact-match Python patch (aborts instead of silently no-op'ing if the file content doesn't match what's expected — the safer choice after the first manual save was lost):

```
cd api/routers
cp journal_router.py journal_router.py.bak

python3 - <<'EOF'
with open("journal_router.py") as f:
    content = f.read()

old_get = '''@router.get("/entries/{entry_id}")
async def get_entry(request: Request, entry_id: str, entry_service: EntryService = Depends(
    get_entry_service)):
    \"\"\"
    TODO: Implement this endpoint to return a single journal entry by ID
    \"\"\"
    raise HTTPException(status_code=501, detail="Not implemented - complete this endpoint
```

```

t!")'''

new_get = '''@router.get("/entries/{entry_id}")
async def get_entry(entry_id: str, entry_service: EntryService = Depends(get_entry_servic
e)):
    \"\"\"Get a single journal entry by ID.\"\"\"
    entry = await entry_service.get_entry(entry_id)
    if not entry:
        raise HTTPException(status_code=404, detail="Entry not found")
    return entry'''

old_delete = '''@router.delete("/entries/{entry_id}")
async def delete_entry(entry_id: str, entry_service: EntryService = Depends(get_entry_serv
ice)):
    \"\"\"
    TODO: Implement this endpoint to delete a specific journal entry
    \"\"\"
    raise HTTPException(status_code=501, detail="Not implemented - complete this endpoin
t!")'''

new_delete = '''@router.delete("/entries/{entry_id}")
async def delete_entry(entry_id: str, entry_service: EntryService = Depends(get_entry_serv
ice)):
    \"\"\"Delete a specific journal entry.\"\"\"
    existing_entry = await entry_service.get_entry(entry_id)
    if not existing_entry:
        raise HTTPException(status_code=404, detail="Entry not found")
    await entry_service.delete_entry(entry_id)
    return {"detail": "Entry deleted successfully"}'''

if old_get not in content or old_delete not in content:
    print("MISMATCH – aborting, no changes made.")
else:
    content = content.replace(old_get, new_get).replace(old_delete, new_delete)
    with open("journal_router.py", "w") as f:
        f.write(content)
    print("Patched successfully.")
EOF

cd ../../

```

Backs up the router file, then does an exact string-match replace of both stub functions. Refuses to write anything if the file doesn't match the expected stub text exactly, rather than risking a partial or corrupted edit.

✓ Final code — `api/routers/journal_router.py`

```

@router.get("/entries/{entry_id}")
async def get_entry(entry_id: str, entry_service: EntryService = Depends(get_entry_servic
e)):
    """Get a single journal entry by ID."""
    entry = await entry_service.get_entry(entry_id)
    if not entry:
        raise HTTPException(status_code=404, detail="Entry not found")
    return entry

@router.delete("/entries/{entry_id}")
async def delete_entry(entry_id: str, entry_service: EntryService = Depends(get_entry_serv

```

```
ice)):
    """Delete a specific journal entry."""
    existing_entry = await entry_service.get_entry(entry_id)
    if not existing_entry:
        raise HTTPException(status_code=404, detail="Entry not found")
    await entry_service.delete_entry(entry_id)
    return {"detail": "Entry deleted successfully"}
```

Both handlers call `entry_service.get_entry()` and `entry_service.delete_entry()`, which were already implemented correctly in `entry_service.py` — the router simply never called them. No changes were needed in any other file.

Verification — after the fix

```
ID=<same-or-new-id>

echo "GET by id:"
curl -s -o /dev/null -w "%{http_code}\n" $BASE/entries/$ID

echo "DELETE by id:"
curl -s -o /dev/null -w "%{http_code}\n" -X DELETE $BASE/entries/$ID

echo "GET by id again (should be gone now):"
curl -s -o /dev/null -w "%{http_code}\n" $BASE/entries/$ID
```



Result: 200 / 200 / 404 — confirmed on the second pass, after re-verifying rather than trusting the first "looks fixed" state. All 6 endpoints now functional.

Phase 3 — Containerize with Docker, Push to ACR

Goal: a tested image sitting in Azure Container Registry.



Command	What it does
<code>docker compose up --build</code>	Builds and runs the app + Postgres together locally
<code>docker run --rm -v /var/run/docker.sock:/var/run/docker.sock aquasec/trivy image learningsteps:local</code>	Scans the local image for known vulnerabilities
<code>az acr create --resource-group <rg> --name <acr-name> --sku Basic</code>	Creates the Azure Container Registry
<code>az acr login --name <acr-name></code>	Authenticates Docker against ACR
<code>docker build -t <acr-name>.azurecr.io/learningsteps:v1 .</code>	Builds the image tagged for ACR
<code>docker push <acr-name>.azurecr.io/learningsteps:v1</code>	Uploads the image to the registry
<code>az acr repository list --name <acr-name> --output table</code>	Confirms the image landed in ACR

Dockerfile:

```
FROM python:3.12-slim AS builder
WORKDIR /app
COPY api/requirements.txt .
RUN pip install --no-cache-dir --user -r requirements.txt

FROM python:3.12-slim
WORKDIR /app
```

```

RUN useradd -m appuser \
  && apt-get update && apt-get upgrade -y && rm -rf /var/lib/apt/lists/*
COPY --from=builder /root/.local /home/appuser/.local
COPY api/ .
ENV PATH=/home/appuser/.local/bin:$PATH
USER appuser
EXPOSE 8000
HEALTHCHECK --interval=30s --timeout=3s CMD curl -f http://localhost:8000/ || exit 1
CMD ["uvicorn", "main:app", "--host", "0.0.0.0", "--port", "8000"]

```

✓ **Checkpoint result:** Trivy scan — 159 total findings (44 HIGH, 0 CRITICAL, rest MEDIUM/LOW/UNKNOWN). 0 CRITICAL passes the project's own gate. The 44 HIGH findings are Debian 13 base-OS packages (util-linux, systemd, glibc, ncurses, perl) with no fix yet published upstream — confirmed unchanged even after `apt-get upgrade`. Image tagged and pushed to ACR as `learningsteps:v1`. Status: ✓ Resolved.

Phase 4 — Infrastructure as Code with Terraform

Goal: the whole environment defined in code, reproducible with `terraform apply`.



Command	What it does
<code>terraform init</code>	Downloads providers, initializes the working directory
<code>terraform plan -out=tfplan</code>	Shows what will change, saves the plan
<code>trivy config .</code>	Scans Terraform for misconfigurations (replaces deprecated tfsec)
<code>terraform apply tfplan</code>	Creates the planned resources in Azure
<code>terraform destroy</code>	Tears everything down (for the recovery demo)

Resources provisioned: VNet + AKS/DB subnets, AKS cluster, PostgreSQL Flexible Server, ACR, Key Vault. Full `.tf` examples are in the attached build guide.

⚠ Three real Azure constraints hit during apply, each fixed in `main.tf`:

- Postgres + VNet injection requires a Private DNS zone** — added `azurerm_private_dns_zone` + `azurerm_private_dns_zone_virtual_network_link`, referenced via `private_dns_zone_id`, and set `public_network_access_enabled = false` (Azure rejects public access alongside VNet integration).
- AKS default service CIDR (10.0.0.0/16) collided with the VNet's own range** — set explicit `service_cidr = "10.1.0.0/16"` and `dns_service_ip = "10.1.0.10"` in `network_profile`.
- Subscription quota: 0 vCPU for Standard_B2s/Bsv2Family in westeurope** — switched the node pool to the `Standard_D2s_v6` family instead, which has separate quota headroom on this subscription, keeping the original `node_count = 2` (D-family and B-family vCPU quotas are tracked independently by Azure).



Final vs. earlier draft — what changed and why

Setting	Earlier draft	Final <code>main.tf</code>	Why changed / benefit
AKS <code>oidc_issuer_enabled</code>	not set (defaults false)	<code>true</code>	Required ahead of Phase 7: GitHub Actions OIDC federated login needs the cluster's OIDC issuer URL to exist — added early so Phase 7 won't need a rebuild
AKS <code>vm_size</code>	<code>Standard_B2s_v2</code>	<code>Standard_D2s_v6</code>	The B-series family had 0 vCPU quota on this subscription in westeurope (confirmed by the apply error). D-series draws from a separate quota pool that had headroom, so 2 full nodes fit without a quota increase request

Setting	Earlier draft	Final <code>main.tf</code>	Why changed / benefit
AKS <code>node_count</code>	<code>1</code> (reduced to fit B-series quota)	<code>2</code>	Restored to 2 once the VM family switch freed enough quota — keeps genuine node-level self-healing demonstrable for Phase 6/9, not just pod-level
Postgres <code>zone</code>	not set	<code>"1"</code>	Pins the server to a specific availability zone instead of leaving placement to Azure's default — more predictable behavior, avoids zone-related placement conflicts with the VNet/subnet

✅ **Checkpoint result:** all 7 resources created and verified live (`az aks show` confirmed `Succeeded` despite a transient polling error in the CLI output). Trivy config misconfigurations reduced 6 → 3 (RBAC enabled, network policy set, Key Vault network ACLs deny-by-default + purge protection fixed). Remaining 3 findings (AKS API IP restriction — CRITICAL, AKS OMS logging — MEDIUM, Key Vault soft-delete retention days — MEDIUM) left open by decision, not part of the core project requirement. Full `terraform destroy` → `apply` reproducibility not yet demonstrated end-to-end (reserved for Phase 9 final demo). Status: ✅ Resolved.

Phase 5 — Connect PostgreSQL + Key Vault

Goal: the app reads its DB connection string securely from Key Vault, not a hardcoded value.



Command	What it does
<code>az postgres flexible-server show --name <db> --resource-group <rg> --query fullyQualifiedDomainName -o tsv</code>	Gets the Postgres FQDN to build the connection string
<code>az keyvault secret set --vault-name <kv> --name database-url --value "..."</code>	Stores the DB connection string as a secret
<code>az keyvault set-policy --name <kv> --upn <your-user> --secret-permissions get list set</code>	Grants your own account permission to manage secrets (needed before <code>secret set</code> will work)
<code>az keyvault network-rule add --name <kv> --ip-address <ip> / --vnet-name <vnet> --subnet <subnet></code>	Allows a specific IP or the AKS subnet through the vault's deny-by-default firewall (Phase 4's <code>network_acls</code>)
<code>az network vnet subnet update --name aks-subnet --service-endpoints Microsoft.KeyVault</code>	Required before a subnet can be added as a Key Vault network rule
<code>az aks enable-addons --addons azure-keyvault-secrets-provider</code>	Enables the Key Vault CSI driver on the cluster
<code>az aks show --query addonProfiles.azureKeyvaultSecretsProvider.identity.objectId</code>	Gets the CSI driver's managed identity object ID
<code>az keyvault set-policy --object-id <csi-identity> --secret-permissions get list</code>	Grants the CSI driver's identity permission to read the secret
<code>az aks update --attach-acr <acr-name></code>	Grants AKS pull permission on the ACR (separate from Key Vault access, needed for the pod's image to actually pull)

SecretProviderClass (`k8s-manifests/secretproviderclass.yaml`):

```

apiVersion: secrets-store.csi.x-k8s.io/v1
kind: SecretProviderClass
metadata:
  name: learningsteps-secrets
spec:
  provider: azure
  parameters:
    usePodIdentity: "false"
    useVMManagedIdentity: "true"

```

```

userAssignedIdentityID: "<csi-driver-client-id>"
keyvaultName: learningsteps-kv-m3p2acr
tenantId: "<tenant-id>"
objects: |
  array:
  - |
    objectName: database-url
    objectType: secret
secretObjects:
- secretName: learningsteps-db-secret
  type: Opaque
  data:
  - objectName: database-url
    key: DATABASE_URL

```



Four real blockers hit and fixed, in order:

Own account lacked `secrets set` permission on the vault — granted via `az keyvault set-policy --upn .`

Key Vault firewall (`default_action = Deny` from Phase 4) blocked both the local CLI call and the pod's mount request — added the local IP as a rule for the CLI step, then added the AKS subnet as a VNet rule (after enabling the `Microsoft.KeyVault` service endpoint on that subnet) so the pod itself could reach the vault.

`SecretProviderClass` was missing identity parameters — CSI mount failed with "failed to get credentials, nodePublishSecretRef secret is not set" until `usePodIdentity`, `useVMManagedIdentity`, and `userAssignedIdentityID` (the CSI driver's own client ID) were added.

Pod stuck in `ImagePullBackOff` — AKS's kubelet identity had no pull rights on ACR; fixed with `az aks update --attach-acr .`




Checkpoint result: pod running, `DATABASE_URL` confirmed present via `kubectrl exec ... env`, and a live `POST /entries` against the Azure Postgres instance returned a created entry with a real UUID — full path (Key Vault → CSI → pod env → app → DB) verified working. Schema (`database_setup.sql`) had to be loaded manually via a temporary in-cluster pod, since Postgres is VNet-only and unreachable directly from Kali. Status:  Resolved.

Phase 6 — Kubernetes Orchestration

Goal: the app self-heals and is reachable through a public IP.

 **Mapping how Helm fits into the Step 6 pipeline**



Replaces the raw `kubectl apply` used for Phase 5 testing — same resources, now managed as one versioned Helm release instead of loose YAML files.

📋 Outlining Helm deployment steps and chart configuration values

#	Step	Why	Command
1	Update chart values	Point the chart at the real ACR image + real Key Vault, matching what Phase 5 proved working	Edit <code>values.yaml</code>
2	Add CSI identity params to chart template	The plain YAML <code>SecretProviderClass</code> needed <code>usePodIdentity</code> / <code>useVMManagedIdentity</code> / <code>userAssignedIdentityID</code> to mount — the chart's template is missing them	Edit <code>templates/secretproviderclas</code>
3	Lint the chart	Catch YAML/template errors before touching the cluster	<code>helm lint ./learningsteps-ch</code>
4	Preview rendered manifests	Confirm the chart produces the same resources as the raw YAML that already works	<code>helm template learningsteps ./learningsteps-chart -f values.yaml</code>
5	Install/adopt the release	A Deployment named <code>learningsteps</code> already exists (Phase 5 raw YAML) — Helm takes it over as a managed release	<code>helm upgrade --install learningsteps ./learningstep -n default</code>
6	Verify release status	Confirms Helm sees it as deployed, not failed	<code>helm status learningsteps</code>
7	Verify pods/service running	Same check as Phase 5, now under Helm's management	<code>kubectl get pods,svc -l app=learningsteps</code>
8	Prove self-healing	Core Phase 6 checkpoint — kill a pod, confirm it's replaced automatically	<code>kubectl delete pod -l app=learningsteps && kubectl pods -w</code>
9	Confirm HPA active	Chart includes autoscaling — verify it's watching the deployment	<code>kubectl get hpa</code>

📄 Required files (before starting)

File	Status	Purpose
<code>learningsteps-chart/Chart.yaml</code>	Already exists (attached)	Chart metadata

File	Status	Purpose
<code>learningsteps-chart/values.yaml</code>	Needs edits below	Image repo/tag, Key Vault name/tenant
<code>learningsteps-chart/templates/deployment.yaml</code>	Already exists (attached)	Pod spec, env, probes
<code>learningsteps-chart/templates/service.yaml</code>	Already exists (attached)	LoadBalancer Service
<code>learningsteps-chart/templates/hpa.yaml</code>	Already exists (attached)	Autoscaling rule
<code>learningsteps-chart/templates/secretproviderclass.yaml</code>	Needs edits below	Key Vault CSI mount — missing identity params

`values.yaml` edits:

```
image:
  repository: m3p2acr.azurecr.io/learningsteps
  tag: v1

keyVault:
  name: learningsteps-kv-m3p2acr
  tenantId: "c0531484-50fa-4ff2-97df-b1f2d4810603"
```

`templates/secretproviderclass.yaml` addition (under `parameters:`):

```
usePodIdentity: "false"
useVMManagedIdentity: "true"
userAssignedIdentityID: "2c0bdb35-7155-4e54-840c-948efead9a3"
```



Command	What it does
<code>az aks get-credentials --name learningsteps-aks --resource-group <rg></code>	Configures kubectl to talk to the cluster
<code>kubectl apply -f k8s-manifests/</code> or <code>helm install learningsteps ./learningsteps-chart</code>	Deploys the app (raw YAML or Helm chart)
<code>kubectl get pods</code> / <code>kubectl get svc learningsteps-svc</code>	Checks pod status and the public IP
<code>kubectl delete pod <name></code>	Kills a pod to prove the Deployment replaces it automatically



Checkpoint result: Helm release `learningsteps` deployed successfully (chart rebuilt from scratch on Kali, raw-YAML Phase 5 resources removed first since Helm refuses to adopt resources without its own ownership labels). Self-healing confirmed — deleted pod replaced automatically with a new name. HPA active and reading real CPU metrics (`2%/70%`). Public LoadBalancer IP assigned. Full chain verified live: `curl` to the public IP → `POST /entries` → `200 OK` with a real created entry, proving public access → Helm-managed pod → Key Vault secret → Postgres all work together. Status: Resolved.



Discovery — Azure auto-stops idle resources on this subscription. After a period of inactivity, both the AKS cluster and the Postgres Flexible Server were found **Stopped**, causing DNS/connection failures that looked like config bugs. Fix each time this happens:

```
az aks start --name learningsteps-aks --resource-group module3-project2
az postgres flexible-server start --resource-group module3-project2 --name learnin
gsteps-db-m3p2acr
```

Check status first with `az aks show ... --query powerState.code` / `az postgres flexible-server show ... --query state` before assuming a real error — this will likely recur between work sessions.

⚡ What Was Missing & Required Changes — Step 6



A focused log of exactly what was absent going into Step 6 and the fixes applied — kept separate from the general Problems log for quick reference.

1. AKS cluster found stopped

```
az aks start --name learningsteps-aks --resource-group module3-project2
```

2. Postgres Flexible Server found stopped

```
az postgres flexible-server start \  
  --resource-group module3-project2 \  
  --name learningsteps-db-m3p2acr
```

3. `learningsteps-chart/` folder did not exist on Kali

It only ever existed as the Notion attachment — never unzipped or recreated on the actual machine. Created from scratch, 7 files:

```
learningsteps-chart/Chart.yaml  
learningsteps-chart/values.yaml  
learningsteps-chart/templates/_helpers.tpl  
learningsteps-chart/templates/deployment.yaml  
learningsteps-chart/templates/service.yaml  
learningsteps-chart/templates/hpa.yaml  
learningsteps-chart/templates/secretproviderclass.yaml
```

4. Helm refused to adopt the Phase 5 raw-YAML resources

`helm upgrade --install` failed: *"Deployment learningsteps ... exists and cannot be imported into the current release"* — Helm requires its own ownership labels (`app.kubernetes.io/managed-by: Helm`, `meta.helm.sh/release-name`, etc.), which raw `kubectl apply` never adds. Fix: delete the Phase 5 resources and let Helm create fresh ones under the same names — zero data loss, since these are stateless app resources, not the database.

```
kubectl delete deployment learningsteps  
kubectl delete service learningsteps-svc --ignore-not-found  
kubectl delete secretproviderclass learningsteps-secrets --ignore-not-found  
helm upgrade --install learningsteps ./learningsteps-chart -n default
```



Standing reminder for every future session: before assuming something is broken, check power state first — `az aks show --query powerState.code` and `az postgres flexible-server show --query state`. Both resources have been found auto-stopped after idle periods on this subscription.

Phase 7 — CI/CD with GitHub Actions + Security Gates

Goal: every push builds, scans, and deploys automatically — a bad finding stops it.



Command	What it does
<code>az ad app create --display-name learningsteps-gh-actions</code>	Creates the Azure AD app used for GitHub OIDC login
<code>az ad app federated-credential create --id <app-id> --parameters '{...}'</code>	Trusts GitHub's OIDC token for this repo/branch — no stored password
<code>az role assignment create --assignee <app-id> --role Contributor --scope ...</code>	Grants the app permission to manage the resource group

Gate logic:

- trufflehog `--fail` → any leaked secret always stops the pipeline (no severity threshold)
- Trivy `exit-code: '1'` on image scan → stops on HIGH/CRITICAL vulnerabilities
- `trivy config` on Terraform → stops on HIGH/CRITICAL misconfigurations (replaces deprecated tfsec)
- Azure login uses OIDC via `azure/login@v2` — the exact image that passed scanning is what gets deployed

Full workflow YAML: see attached build guide.



Checkpoint result: full pipeline green end to end — secrets scan → lint/test → build → Trivy image scan → OIDC login → ACR push → Helm deploy. New pod deployed live by the pipeline, public endpoint re-confirmed `200`. Push-to-auto-deploy proven. Still open: deliberately break the pipeline (bad package or secret) to prove the gate blocks it — reserved for Phase 9.



Re-confirmed after bumping `azure/login@v2` → `@v3` — pushed from Kali, full re-run succeeded (run #9, 1m 25s): secrets scan → lint/test → build → Trivy scan → Azure login (OIDC) → ACR push → Helm deploy, all green. Confirms the pipeline is stable across an action-version bump, not just a one-off pass.



Six real issues hit setting this up, in order:

Original repo is not owned — `CyberstepsDE/learningsteps` can't be pushed to directly. Forked to `Hany-Dh/learningsteps` and repointed the local remote (`git remote set-url origin ...`) before any pipeline work.

GitHub password auth removed — `git push` needs a Personal Access Token, not the account password.

trufflehog step failed: "flag 'fail' cannot be repeated" — the action already passes `--fail` by default; the added `extra_args: --fail` duplicated it. Fix: remove `extra_args` entirely.

Trivy gate failed the build at 44 HIGH / 0 CRITICAL — same baseline accepted back in Phase 3 (unfixed upstream Debian CVEs). Narrowed both Trivy steps to `severity: CRITICAL` to match the project's actual pass bar, since 0 CRITICAL is what has been used project-wide.

Terraform scan step failed: `infra-terraform` not found — that folder lives outside the `learningsteps` git repo on Kali (separate directory), so the workflow's repo checkout never contains it. Removed the "Scan Terraform" step from this workflow rather than merging directory structures.

OIDC login failed: "No matching federated identity record" — GitHub's actual token subject for a forked repo is `repo:Hany-Dh@106172590/learningsteps@<run-id>:ref:refs/heads/main` (an `@`-separated owner ID), not the plain `repo:Hany-Dh/learningsteps:ref:refs/heads/main` form first registered. Deleted and recreated the federated credential with the exact subject from the error log.

Helm deploy failed: `path "./learningsteps-chart" not found` — the chart directory existed locally under `learningsteps-project/` but was never copied into the `learningsteps/` repo folder that actually gets pushed. Copied it in and committed.

Phase 8 — Optional: Monitoring

Goal: live dashboard showing request volume, latency, and DB health.



Command	What it does
<code>pip install prometheus-client prometheus-fastapi-instrumentator</code>	Adds metrics support to the FastAPI app
<code>helm install monitoring prometheus-community/kube-prometheus-stack --namespace monitoring --create-namespace</code>	Deploys Prometheus + Grafana to the cluster
<code>kubectll port-forward -n monitoring svc/monitoring-grafana 3000:80</code>	Opens Grafana locally to build the dashboard

✓ **Checkpoint:** real traffic against the API produces visible dashboard changes.

|| **Status: paused by decision** — this step is marked optional in the project brief, so work stopped here to prioritize Step 9's required proofs.

Completed:

App instrumented (`prometheus-client` + `prometheus-fastapi-instrumentator`), `/metrics` live with real data.

Prometheus + Grafana deployed via Helm (slimmed install).

Service named port + labels added for discovery; `ServiceMonitor` created and scrape target confirmed `up` with no errors.

Grafana login working; Prometheus data source confirmed present and set as default.

Remaining:

Prebuilt FastAPI dashboard (ID 16110) blocked — no Loki data source installed, import stays disabled.

Manual dashboard: 3 panels planned (HTTP Request Volume, API Latency p95, App/DB Reachability) — panel 1 query entered, not yet confirmed rendering or saved.

Live traffic test against the dashboard not yet run.

Estimated remaining effort: ~10–15 minutes if resumed.

📏 Resource Sizing Decision — Full vs. Slim Install

📌 AKS runs 2× `Standard_D2s_v6` (2 vCPU / 8 GiB each, 4 vCPU / 16 GiB total). A default kube-prometheus-stack install carries components beyond what this project's checkpoint needs, so a slimmed install was chosen instead of the defaults.

Component	Full default	Slimmed install	Required by project?
Prometheus server	~500m CPU / ~2Gi RAM	200m CPU / 512Mi RAM	✓ Yes — explicitly required (metrics collection)
Grafana	~100m CPU / ~256Mi RAM	50m CPU / 128Mi RAM	✓ Yes — explicitly required (dashboard)
Alertmanager	~100m CPU / ~256Mi RAM, full alerting pipeline	Disabled	✗ No — project asks for dashboards, not alert routing
kube-state-metrics	~100m CPU / ~128Mi RAM	Included, unchanged	✓ Yes, indirectly — feeds Grafana's default dashboards
node-exporter	Small DaemonSet × 2 nodes	Included, unchanged	⚠ Optional — required metrics are app-level (HTTP volume, latency, DB health), not node-level
Persistent storage (PVCs)	Often enabled by default	Not set — ephemeral	✗ No — demo only needs live dashboard changes, not retained history
Total requested (approx.)	~1 vCPU / ~3Gi RAM	~350–450m CPU / ~900Mi RAM	—

Safe to exclude without affecting the checkpoint: Alertmanager, PVCs, and optionally node-exporter. Required to stay: Prometheus, Grafana, kube-state-metrics — the three that the checkpoint actually depends on.



Capacity check before install: `kubecttl top nodes` showed real usage at 7% / 4% CPU and ~25% memory on both nodes — ample headroom for the slim install. Reserved (requested) resources sat higher (50% / 24% CPU) but that reflects reservations, not actual load. One node showed CPU/memory **limits** overcommitted on paper (536% / 185%) — not a blocker, since limits only matter if multiple pods spike CPU simultaneously, but worth monitoring.

#	Item	Status
1	FastAPI app instrumented (<code>prometheus-client</code> + <code>prometheus-fastapi-instrumentator</code>)	✅ Done
2	<code>/metrics</code> endpoint live and exposing real data	✅ Done
3	Prometheus + Grafana deployed via Helm (slimmed install)	✅ Done
4	Service named port + labels added for discovery	✅ Done
5	<code>ServiceMonitor</code> created and scraping confirmed <code>up</code>	✅ Done
6	Grafana login working	✅ Done
7	Prometheus data source confirmed present in Grafana	✅ Done
8	Prebuilt FastAPI dashboard import	❌ Blocked — missing Loki data source, abandoned
9	Manual dashboard panel 1 — HTTP Request Volume	🔄 In progress — query entered, not yet confirmed rendering
10	Manual dashboard panel 2 — API Latency (p95)	🟪 Not started
11	Manual dashboard panel 3 — App/DB Reachability	🟪 Not started
12	Save the dashboard	🟪 Not started
13	Generate traffic and confirm dashboard moves live (the actual checkpoint)	🟪 Not started

Bottom line: the entire pipeline (app → Prometheus) is proven working — items 1–7 are the hard infrastructure part and it's done. What's left (8–13) is just Grafana UI clicking to visualize what's already flowing



Feature Addition — Dual Content-Type Support (JSON + text/plain)



Added before Step 9 at user request: `POST /entries` and `PATCH /entries/{id}` previously accepted `application/json` only. Both now accept **either** `application/json` **or** `text/plain`, routed through the same `EntryService` /PostgreSQL logic either way.

Files changed: `api/routers/journal_router.py`, `api/models/entry.py`

`entry.py` — `struggle` and `intention` made `Optional` on `EntryCreate` (plain text only populates `work`):

```
struggle: Optional[str] = Field(default=None, max_length=256, ...)
intention: Optional[str] = Field(default=None, max_length=256, ...)
```

`journal_router.py` — a shared `parse_request_body()` helper reads `Content-Type` and either parses JSON or wraps raw plain text into `{"work": text}`; both `POST` and `PATCH` call it. `openapi_extra` was added to both route decorators so Swagger (`/docs`) actually lists both content types as selectable — without it, FastAPI can't generate a schema for a raw `Request` body and `/docs` silently shows JSON only, which is what first looked like a bug.

⚠️ Two iterations needed: the first drafted router was missing the `get_entry_service` dependency and the other four endpoints (GET all, GET one, DELETE one, DELETE all) — caught by comparing against the actual original file rather than a memory-reconstructed version. The corrected final version preserves 100% of the original file's comments, TODOs, and structure, with only the two target endpoints modified.

🔧 Verification — all 4 combinations tested live

```
# JSON POST
curl -X POST http://4.175.51.84/entries -H "Content-Type: application/json" -d '{"work":"json test 2"}'

# text/plain POST
curl -X POST http://4.175.51.84/entries -H "Content-Type: text/plain" --data-binary "Learn ed about openapi_extra today."

# JSON PATCH
curl -X PATCH http://4.175.51.84/entries/<id> -H "Content-Type: application/json" -d '{"work":"updated via json"}'

# text/plain PATCH
curl -X PATCH http://4.175.51.84/entries/<id> -H "Content-Type: text/plain" --data-binary "Updated via plain text."
```

✅ **Checkpoint result:** all 4 combinations confirmed working against live Azure Postgres. `openapi.json` confirms both content types registered: `['application/json', 'text/plain']`. Swagger UI (`/docs`) now shows both options for POST and PATCH. Deployed via the Step 7 pipeline (push → auto-deploy), no manual kubectl steps needed. Status: ✅ Resolved.

Phase 9 — Final Demonstration



Demo	What to show
Deployment	Push a small change, show it live on the public IP automatically
Security enforcement	Introduce a vulnerable package or insecure Terraform rule, show the pipeline blocking it
Infrastructure recovery	<code>terraform destroy</code> then <code>terraform apply</code> rebuilds everything
Documentation	Submit app code, Dockerfile, Terraform, K8s manifests, workflows, README

📌 Planned proof steps below. Project confirmed functionally complete and ready for presentation — these two demos are **documented as a plan, not executed**, by decision; running them (especially infra recovery) would tear down and rebuild the live demo environment.

🛡️ Security Enforcement — prove the pipeline blocks a bad push

#	Step	Why	Command / Code
1	Create a test branch	Keeps <code>main</code> clean, easy to revert	<code>git checkout -b security-test</code>
2	Add a known-vulnerable Python package	Triggers a real CRITICAL finding from Trivy's image scan	Add <code>pyyaml==5.3.1</code> to <code>api/requirements.txt</code>
3	Commit and push	Fires the pipeline	<code>git add api/requirements.txt && git commit -m "Test: introduce vulnerable dependency" && git push origin security-test</code>
4	Merge/push to <code>main</code>	Pipeline only triggers on push to <code>main</code>	<code>git push origin security-test:main</code>
5	Watch GitHub Actions	Confirm Trivy step fails and pipeline stops before deploy	github.com/Hany-Dh/learningsteps/actions
6	Capture evidence	Needed for demo/writeup	Screenshot the red step + CRITICAL finding
7	Revert the change	Pipeline shouldn't stay broken	<code>git revert HEAD && git push origin main</code>
8	Confirm pipeline passes again	Proves the gate works both ways	Watch Actions tab for green run

Infrastructure Recovery — terraform destroy → apply

#	Step	Why	Command / Code
1	Confirm state matches reality	Avoid destroying with stale state	<code>terraform plan -var="db_admin_password=<pwd>"</code> (expect no changes)
2	Note current outputs	Reference values for comparison after rebuild	<code>terraform output</code>
3	Destroy everything	Required proof per project brief	<code>terraform destroy -var="db_admin_password=<pwd>"</code>
4	Confirm resource group empty	Verifies destroy actually worked	<code>az resource list --resource-group module3-project2 -o table</code>
5	Re-apply from code only	The actual recovery proof	<code>terraform apply -var="db_admin_password=<pwd>"</code>
6	Re-attach ACR pull permission	Not in Terraform state — was set via <code>az aks update</code> separately	<code>az aks update --name learningsteps-aks --resource-group module3-project2 --attach-acr m3p2acr</code>
7	Re-create Key Vault secret + network rules	Set via <code>az</code> , not Terraform	Re-run Phase 5's <code>az keyvault secret set</code> • <code>network-rule add</code>
8	Re-deploy the app	Infra exists again but nothing is running yet	<code>helm upgrade --install learningsteps ./learningsteps-chart -n default</code>
9	Verify public endpoint live	Final proof the full stack recovered	<code>curl -s -o /dev/null -w "%{http_code}\n" http://<new-ip>/docs</code>
10	Capture evidence	Needed for demo/writeup	Screenshot destroy completion, apply completion, working endpoint



Note: destroy wipes the current public IP and any Postgres data not otherwise backed up — why this was documented as a plan rather than executed against the live demo environment.



Decision: the two proof procedures below are planned and documented, but **not executed** — project considered functionally complete and finalized for presentation as of this point. Running `terraform destroy` would wipe the live public IP and current Postgres data; the security-enforcement test would require reverting a deliberately-broken dependency. Both are low-risk, well-understood procedures that can be run live during an actual demo/presentation if required, rather than pre-emptively on the working environment.



Problems & Issues Discussed



Running log of real issues hit while working through this project, and how they were resolved.

Issue	Cause	Resolution
<code>az</code> install failed: "Unable to find a package for your system"	Kali is a rolling release; Microsoft's Azure CLI apt repo has no package matching <code>kali-rolling</code>	Installed via <code>pipx install az</code> apt — sidesteps distro-coder entirely
<code>tfsec: command not found</code> after running the install script	tfsec was deprecated and merged into Trivy by Aqua Security; its install script is effectively dead	Dropped tfsec entirely — Ter uses <code>trivy config <dir></code> for and the GitHub Actions step
<code>trivy config infra-terraform/</code> returned a FATAL "no such file or directory"	Expected — <code>infra-terraform/</code> doesn't exist yet since Terraform work (Phase 4) hasn't started	Not a real error; the checks-k succeeded. Re-run once the exists
<code>docker ps</code> failed after restart even though the docker service was active	Shell session didn't have the <code>docker</code> group applied	<code>sudo usermod -aG docker \$USER</code> , then <code>docker run --r</code> succeeded. (Earlier, a leftover podman-docker also caused fix with <code>unset DOCKER_HOST</code>)
Stub-endpoint fix appeared applied but GET/DELETE by id still returned 501 after "fixing"	The router edit was never actually saved to <code>journal_router.py</code> — marked resolved before the verification round actually confirmed it	Re-applied the fix with a Pyth (exact string match, aborts or of silently failing), confirmed up the change, re-ran GET/D 200/200/404 this time
<code>docker compose up --build</code> failed: "open Dockerfile: no such file or directory"	File was saved as <code>dockerfile.txt</code> instead of <code>Dockerfile</code>	<code>mv dockerfile.txt Dockerfi</code>
Postgres container failed to bind port 5432	A native PostgreSQL service was already installed on Kali (separate from this project, listening on <code>:::1:5432</code>), plus the leftover Step 2 standalone container	Stopped/removed the Step 2 remapped compose's Postgr <code>5433:5432</code> , left the native Ka untouched — internal contain traffic uses <code>postgres:5432</code> mapping
App container failed to bind port 8000	Step 2's <code>uvicorn --reload</code> process was still running on the host	<code>sudo fuser -k 8000/tcp</code> , the host
App container ran but had no port mapping in <code>docker compose ps</code>	Container was created during an earlier failed attempt and reused without rebinding the port	<code>docker compose down && doc build -d --force-recreate</code>
<code>POST /entries</code> returned raw "Internal Server Error" inside the container	<code>.env</code> 's <code>DATABASE_URL</code> still pointed at <code>localhost:5433</code> (a Step 2 leftover) — unreachable from inside the app container, which must use the compose service name	Changed <code>.env</code> to <code>postgres: service name/port, independ 5433</code> mapping), recreated th
<code>az acr create</code> rejected the registry name	Name contained a dash (<code>M3P2-ACR</code>) — ACR names must be alphanumeric only	Renamed to <code>m3p2acr</code>
<code>az acr login</code> / <code>docker push</code> hung or failed with DNS/auth errors	<code>az group create</code> was never run — the resource group didn't exist, so <code>az acr create</code> silently never fully provisioned	Ran <code>az group create --name -location westeurope</code> first, t commands succeeded
<code>terraform apply</code> failed: resource group / ACR / VNet "already exists" errors	These resources were created manually via <code>az</code> in Step 3, so Terraform had no record of them in its state file	<code>terraform import</code> for each (ACR, VNet) before re-running
<code>terraform apply</code> failed: "polling after CreateOrUpdate: unimplemented polling status Unknown" (hit twice — VNet, then AKS)	Transient Azure API polling glitch, not a real failure — resource was	Verified via <code>az ... show --q provisioningState</code> (returned re-ran <code>terraform apply / imp</code>

Issue	Cause	Resolution
	actually created successfully both times	
<code>terraform apply</code> failed: Postgres <code>EmptyPrivateDnsZoneArmResourceId</code>	Azure requires a Private DNS zone when a Postgres Flexible Server uses VNet injection (<code>delegated_subnet_id</code>)	Added <code>azurerm_private_dns</code> , <code>azurerm_private_dns_zone_v</code> resources, linked via <code>private</code>
<code>terraform apply</code> failed: AKS "VM size Standard_B2s is not allowed in your subscription in location westeurope"	Subscription's allowed SKU list doesn't include the original <code>v1</code> size class	Switched to <code>Standard_B2s_v2</code> allowed list
<code>terraform apply</code> failed: Postgres "ConflictingPublicNetworkAccessAndVirtualNetworkConfiguration"	Azure rejects public network access enabled alongside VNet-integrated Postgres	Set <code>public_network_access_e</code> the Postgres resource
<code>terraform apply</code> failed: AKS "ServiceCidrOverlapExistingSubnetsCidr"	Azure CNI's default service CIDR (10.0.0.0/16) collided with the VNet's own address space (also 10.0.0.0/16)	Set explicit non-overlapping <code>"10.1.0.0/16"</code> and <code>dns_serv</code> <code>"10.1.0.10"</code> in AKS <code>network</code>
<code>terraform apply</code> failed: AKS "Insufficient vcpu quota requested 4, remaining 0 for family standardBsv2Family"	Subscription's vCPU quota for that VM family in westeurope is 0 — 2 nodes × 2 vCPU exceeded it	Dropped <code>node_count</code> from 2 vCPU request to fit the quota no node-level HA demo)
<code>git push</code> to the original project repo was about to happen	<code>origin</code> still pointed at <code>CyberstepsDE/learningsteps</code> , not an owned repo	Forked to <code>Hany-Dh/learnings set-url origin https://gith Dh/learningsteps.git</code>
<code>git push</code> failed: "Invalid username or token. Password authentication is not supported"	GitHub removed password auth for git operations	Generated a Personal Access <code>repo</code> scope), used it as the prompt
GitHub Actions: truffleshog step failed — "flag 'fail' cannot be repeated"	The action already passes <code>--fail</code> internally; <code>extra_args: --fail</code> duplicated the flag	Removed <code>extra_args: --fail</code> step entirely
GitHub Actions: Trivy image scan failed the pipeline at 44 HIGH / 0 CRITICAL	Workflow was gated on <code>HIGH,CRITICAL</code> , but these are the same unfixed-upstream findings already accepted as a pass in Phase 3	Changed <code>severity: HIGH,CRITICAL</code> on both Trivy steps project's established pass ba
GitHub Actions: Terraform scan step failed — "infra-terraform ... no such file or directory"	<code>infra-terraform/</code> lives outside the <code>learningsteps</code> git repo on Kali, so it's absent from the Actions checkout	Removed the "Scan Terraform <code>cd.yml</code> rather than merging s structures
GitHub Actions: Azure login (OIDC) failed — "No matching federated identity record"	GitHub's actual OIDC subject for this fork used <code>repo:Hany-Dh@106172590/learningsteps@<run-id>:ref:refs/heads/main</code> , with an <code>@</code> -separated owner ID — the federated credential was registered with a plain <code>Hany-Dh/learningsteps</code> subject that didn't match	Deleted and recreated the fed using the exact subject string error log
GitHub Actions: AKS unreachable during "Deploy with Helm" — DNS lookup failure	AKS cluster was auto-stopped from inactivity (same behavior seen in Phase 6)	<code>az aks start --name learni resource-group module3-proj</code> retriggered the pipeline
GitHub Actions: Helm deploy failed — "path ./learningsteps-chart not found"	Chart existed on Kali under <code>learningsteps-project/learningsteps-chart/</code> , outside the <code>learningsteps/</code> repo that Actions checks out	<code>cp -r learningsteps-chart</code> committed and pushed



Attached Documents (latest versions)

These are regenerated as the project evolves — always check this page for the current version rather than an old download.

[learningsteps-devops-guide.md](#)

Full build guide (Phases 2–9, commands, YAML, Dockerfile)

[learningsteps-setup.txt](#)

Resource setup script (Terraform → ACR → AKS → Key Vault → OIDC, sectioned)

[install-tools-kali.txt](#)

Kali Linux tool install script (idempotent, pipx-based Azure CLI)

[learningsteps-helm-chart.md](#)

Helm chart (Chart.yaml, values.yaml, templates)